

CSCI E-89B Introduction to Natural Language Processing

Harvard Extension School

Dmitry Kurochkin

Fall 2024
Lecture 5

Contents

1 Term Frequency-Inverse Document Frequency (TF-IDF)

- Introduction to TF-IDF
- TF-IDF Example
- Advantages of TF-IDF over BoW
- Modified Versions of TF-IDF
- Hands-On: TF-IDF in Python Using Sklearn

2 Word Embeddings

- Introduction to Word Embeddings
- Limitations of Word Embeddings
- Word2Vec
- GloVe

3 Practical Application in Python

Contents

1 Term Frequency-Inverse Document Frequency (TF-IDF)

- Introduction to TF-IDF
- TF-IDF Example
- Advantages of TF-IDF over BoW
- Modified Versions of TF-IDF
- Hands-On: TF-IDF in Python Using Sklearn

2 Word Embeddings

- Introduction to Word Embeddings
- Limitations of Word Embeddings
- Word2Vec
- GloVe

3 Practical Application in Python

Introduction to TF-IDF

- **Term Frequency-Inverse Document Frequency (TF-IDF)**

- ▶ TF-IDF is a statistical measure used to evaluate the importance of terms in documents relative to a corpus.

- ★ **Term Frequency (TF):** Measures how often a term appears in a document. In simplest form, TF is calculated as:

$$\text{TF}(\text{term}, \text{doc}) = \frac{\text{Number of times term appears in the document}}{\text{Total number of terms in the document}}$$

- ★ **Inverse Document Frequency (IDF):** Assesses how much information a word provides by downscaling terms that occur frequently across the corpus:

$$\text{IDF}(\text{term}) = \ln \left(\frac{\text{Total number of documents}}{\text{Number of documents containing the term}} \right)$$

- ▶ **TF-IDF:**

$$\text{TF-IDF}(\text{term}, \text{doc}) = \text{TF}(\text{term}, \text{doc}) \times \text{IDF}(\text{term})$$

Contents

1 Term Frequency-Inverse Document Frequency (TF-IDF)

- Introduction to TF-IDF
- **TF-IDF Example**
- Advantages of TF-IDF over BoW
- Modified Versions of TF-IDF
- Hands-On: TF-IDF in Python Using Sklearn

2 Word Embeddings

- Introduction to Word Embeddings
- Limitations of Word Embeddings
- Word2Vec
- GloVe

3 Practical Application in Python

TF-IDF Example: Manual Computation

- **Documents:**

- ① Doc 1: "cat dog cat"
- ② Doc 2: "dog mouse dog"
- ③ Doc 3: "dog mouse"
- ④ Doc 4: "mouse cat dog"

- **Vocabulary:** {"cat", "dog", "mouse"}

- **Term Frequency (TF):**

- ▶ $TF(\text{cat})$: Doc 1 = $2/3$, Doc 2 = $0/3$, Doc 3 = $0/2$, Doc 4 = $1/3$
- ▶ $TF(\text{dog})$: Doc 1 = $1/3$, Doc 2 = $2/3$, Doc 3 = $1/2$, Doc 4 = $1/3$
- ▶ $TF(\text{mouse})$: Doc 1 = $0/3$, Doc 2 = $1/3$, Doc 3 = $1/2$, Doc 4 = $1/3$

- **Inverse Document Frequency (IDF):**

$$IDF(\text{cat}) = \ln(4/2) = 0.693$$

$$IDF(\text{dog}) = \ln(4/4) = 0$$

$$IDF(\text{mouse}) = \ln(4/3) = 0.288$$

TF-IDF Example: Manual Computation (Continued)

• TF-IDF Values:

1 Doc 1:

★ **cat**: $2/3 \times \ln(4/2) = 2/3 \times 0.693 \approx 0.462$

★ **dog**: $1/3 \times \ln(4/4) = 1/3 \times 0 = 0$

★ **mouse**: $0/3 \times \ln(4/3) = 0/3 \times 0.288 = 0$

2 Doc 2:

★ **cat**: $0/3 \times \ln(4/2) = 0/3 \times 0.693 = 0$

★ **dog**: $2/3 \times \ln(4/4) = 2/3 \times 0 = 0$

★ **mouse**: $1/3 \times \ln(4/3) = 1/3 \times 0.288 \approx 0.096$

3 Doc 3:

★ **cat**: $0/2 \times \ln(4/2) = 0/2 \times 0.693 = 0$

★ **dog**: $1/2 \times \ln(4/4) = 1/2 \times 0 = 0$

★ **mouse**: $1/2 \times \ln(4/3) = 1/2 \times 0.288 \approx 0.144$

4 Doc 4:

★ **cat**: $1/3 \times \ln(4/2) = 1/3 \times 0.693 \approx 0.231$

★ **dog**: $1/3 \times \ln(4/4) = 1/3 \times 0 = 0$

★ **mouse**: $1/3 \times \ln(4/3) = 1/3 \times 0.288 \approx 0.096$

Contents

1 Term Frequency-Inverse Document Frequency (TF-IDF)

- Introduction to TF-IDF
- TF-IDF Example
- **Advantages of TF-IDF over BoW**
- Modified Versions of TF-IDF
- Hands-On: TF-IDF in Python Using Sklearn

2 Word Embeddings

- Introduction to Word Embeddings
- Limitations of Word Embeddings
- Word2Vec
- GloVe

3 Practical Application in Python

Advantages of TF-IDF over BoW

- Enhances Bag of Words (BoW) by incorporating global term significance.
 - ▶ **Limitations of BoW:** Simply counts occurrences, potentially leading common words to overshadow meaningful terms.
 - ▶ **Advantage of TF-IDF:** Adjusts term weights based on their occurrence in the corpus, offering a more comprehensive measure of term importance.
- Balances term frequency with significance across a corpus.
 - ▶ **Balancing Act:** Combines local document frequency with global significance, refining how insights are drawn.
 - ▶ **Applications:** Widely used for keyword extraction, document classification, and enhancing search engine results.
 - ▶ **Normalization:** Often employed to make TF-IDF scores comparable across documents.

Contents

1 Term Frequency-Inverse Document Frequency (TF-IDF)

- Introduction to TF-IDF
- TF-IDF Example
- Advantages of TF-IDF over BoW
- **Modified Versions of TF-IDF**
- Hands-On: TF-IDF in Python Using Sklearn

2 Word Embeddings

- Introduction to Word Embeddings
- Limitations of Word Embeddings
- Word2Vec
- GloVe

3 Practical Application in Python

Modified Versions of Term Frequency

- **Raw Term Count:**

- ▶ $f(t, d)$ represents the raw count of term t occurrences in document d .
- ▶ Can be affected by document length.

- **Term Frequency (TF):**

- ▶ $TF(t, d) = \frac{f(t, d)}{N_d}$
- ▶ N_d is the total number of terms in document d .
- ▶ Default method in Scikit-learn's 'TfidfVectorizer'.

- **Logarithmically Scaled TF:**

- ▶ $TF_{\ln}(t, d) = 1 + \ln(f(t, d))$ if $f(t, d) > 0$; else 0.
- ▶ Reduces the impact of high frequency terms.

- **Double Normalization (K):**

- ▶ $TF_{\text{norm}}(t, d) = K + (1 - K) \times \frac{f(t, d)}{\max\{f(w, d): w \in d\}}$
- ▶ Balances term frequencies relative to the maximum term frequency in the document.
- ▶ K is typically 0.5.

- **Boolean Term Frequency:**

- ▶ $TF_{\text{bool}}(t, d) = 1$ if the term appears in the document, otherwise 0.
- ▶ Ignores actual frequency, focusing on presence.

Modified Versions of Inverse Document Frequency (IDF)

- **Standard IDF:**

- ▶ $IDF(t) = \ln \left(\frac{N}{DF(t)} \right)$, where $DF(t)$ is the document frequency of term t .
- ▶ High IDF for rare terms across documents.

- **Smoothed IDF:**

- ▶ $IDF_s(t) = \ln \left(\frac{1+N}{1+DF(t)} \right) + 1$
- ▶ Default method in Scikit-learn's 'TfidfVectorizer'.
- ▶ Includes smoothing to prevent division by zero.

- **Probabilistic IDF (ProbIDF):**

- ▶ $IDF_{prob}(t) = \ln \left(\frac{N-DF(t)}{DF(t)} \right)$
- ▶ Considers the probability of document exclusion.

- **Maximal IDF:**

- ▶ $IDF_{max}(t) = \ln \left(\frac{\max(DF)}{1+DF(t)} \right)$
- ▶ Uses maximum document presence as a benchmark.

Contents

1 Term Frequency-Inverse Document Frequency (TF-IDF)

- Introduction to TF-IDF
- TF-IDF Example
- Advantages of TF-IDF over BoW
- Modified Versions of TF-IDF
- Hands-On: TF-IDF in Python Using Sklearn

2 Word Embeddings

- Introduction to Word Embeddings
- Limitations of Word Embeddings
- Word2Vec
- GloVe

3 Practical Application in Python

Hands-On: TF-IDF in Python Using Sklearn

```
from sklearn.feature_extraction.text import TfidfVectorizer

# Example text
documents = ["cat dog cat",
             "dog mouse dog",
             "cat mouse",
             "mouse cat dog"]

# Initialize TfidfVectorizer
vectorizer = TfidfVectorizer()

# Fit and Transform the documents
tfidf_matrix = vectorizer.fit_transform(documents)

# Display the Vocabulary and TF-IDF Representation
print("Vocabulary:\n", vectorizer.get_feature_names_out(), "\n")
print("TF-IDF Representation:\n", tfidf_matrix.toarray())
```

Vocabulary:

['cat', 'dog', 'mouse']

TF-IDF Representation:

```
[[0.81649658 0.40824829 0.          ]
 [0.          0.81649658 0.40824829]
 [0.70710678 0.          0.70710678]
 [0.40824829 0.40824829 0.40824829]]
```

Contents

1 Term Frequency-Inverse Document Frequency (TF-IDF)

- Introduction to TF-IDF
- TF-IDF Example
- Advantages of TF-IDF over BoW
- Modified Versions of TF-IDF
- Hands-On: TF-IDF in Python Using Sklearn

2 Word Embeddings

- Introduction to Word Embeddings
- Limitations of Word Embeddings
- Word2Vec
- GloVe

3 Practical Application in Python

Introduction to Word Embeddings

- **Concept:**

- ▶ Word embeddings are dense vector representations of words.
- ▶ They encode words into numerical vectors where semantically similar words have similar representation.
- ▶ Unlike traditional Bag of Words, embeddings capture context by considering the proximity of words in text.

- **Goal:**

- ▶ Aim to create a high-dimensional space where relationships and meanings between words reflect their context.
- ▶ Facilitate understanding of words in multiple dimensions, addressing ambiguity through context.
- ▶ Enable machines to understand and process text in ways similar to human cognition.


Introduction to Word Embeddings

Word embeddings can be considered an alternative to one-hot encoding:



One-hot word vectors:

- Sparse
- High-dimensional
- Hardcoded




Word embeddings:


- Dense
- Lower-dimensional
- Learned from data

Introduction to Word Embeddings

Example:



Male-Female



Verb tense

Applications of Word Embeddings

- ➊ **Sentiment Analysis:** Improve accuracy by understanding context and nuances in opinions.
- ➋ **Machine Translation:** Capture word equivalencies across languages, enhancing translation quality.
- ➌ **Information Retrieval and Search:** Enable intuitive and relevant search results by understanding related terms and concepts.
- ➍ **Recommendation Systems:** Use vectors to determine item and user similarities, improving recommendations.
- ➎ **Word Similarity and Analogy Tasks:** Enable finding similar or related words, and solve analogy tasks (e.g., "king" is to "queen" as "man" is to "woman").
- ➏ **Text Classification:** Enhance accuracy by providing semantically rich word representations.
- ➐ **Named Entity Recognition (NER):** Facilitate identifying and classifying named entities like people, organizations, and locations.

Applications of Word Embeddings (Continued)

- 8 **Topic Modeling and Clustering:** Assist in grouping similar documents to uncover underlying themes within datasets.
- 9 **Question Answering Systems:** Improve retrieval and understanding of relevant information by capturing word relationships.
- 10 **Language Modeling and Generation:** Enhance sequence probability prediction, aiding in text understanding and generation.
- 11 **Relation Extraction:** Enable the extraction of semantic relationships between entities, enhancing databases and knowledge graphs.
- 12 **Social Media Monitoring:** Detect sentiment trends and conduct topic analysis in social media text.
- 13 **Spelling Correction and Typing Suggestions:** Provide more accurate recommendations by understanding context.

Contents

1 Term Frequency-Inverse Document Frequency (TF-IDF)

- Introduction to TF-IDF
- TF-IDF Example
- Advantages of TF-IDF over BoW
- Modified Versions of TF-IDF
- Hands-On: TF-IDF in Python Using Sklearn

2 Word Embeddings

- Introduction to Word Embeddings
- **Limitations of Word Embeddings**
- Word2Vec
- GloVe

3 Practical Application in Python

Limitations of Word Embeddings

- **Out-of-Vocabulary (OOV) Words:**

- ▶ Static embeddings cannot handle words not seen during training.
- ▶ Subword approaches like FastText can mitigate this.

- **Context Insensitivity (Static Embeddings):**

- ▶ Single vector representation per word fails to capture different meanings in diverse contexts.
- ▶ Modern context-specific embeddings address this limitation.

- **Bias Reflection and Amplification:**

- ▶ Embeddings can encapsulate biases present in training data.
- ▶ These biases can influence downstream applications, requiring debiasing techniques.

- **Resource Intensive:**

- ▶ Training and fine-tuning embeddings require significant computational resources.
- ▶ Infeasible for real-time deployment in limited-resource environments.

Limitations of Word Embeddings (Continued)

- **Semantic Drift:**

- ▶ Static embeddings do not adapt to changes in language over time.
- ▶ Requires periodic retraining to maintain relevance with language evolution.

- **Low-Resource Language Limitations:**

- ▶ Less effective for languages with limited training data, potentially yielding poorer performance.

Contents

1 Term Frequency-Inverse Document Frequency (TF-IDF)

- Introduction to TF-IDF
- TF-IDF Example
- Advantages of TF-IDF over BoW
- Modified Versions of TF-IDF
- Hands-On: TF-IDF in Python Using Sklearn

2 Word Embeddings

- Introduction to Word Embeddings
- Limitations of Word Embeddings
- Word2Vec
- GloVe

3 Practical Application in Python

Word2Vec: Introduction and Key Models

- **Developed by Google:**

- ▶ Created by a team at Google led by Tomas Mikolov in 2013.
- ▶ Aimed to efficiently process vast amounts of text to produce high-quality word embeddings.
- ▶ Revolutionized NLP by significantly improving the computational efficiency of training word vectors.

- **Key Models:**

- ▶ **Skip-gram:**

- ★ Predicts surrounding context words for a given target word.
- ★ Works well with small datasets.
- ★ Effective for modeling infrequent words by leveraging context.

- ▶ **CBOW (Continuous Bag of Words):**

- ★ Predicts a target word based on a given context of surrounding words.
- ★ More computationally efficient than Skip-gram for large datasets.
- ★ Tends to perform better with frequent words.

Word2Vec: CBOW and Applications

- **Properties:**

- ▶ Utilizes shallow neural networks with one hidden layer.
- ▶ Trained using techniques like negative sampling or hierarchical softmax to optimize learning.
- ▶ Capable of capturing semantic relationships, enabling vector arithmetic (e.g., "king" - "man" + "woman" = "queen").

- **Applications:**

- ▶ Enhances document clustering, sentiment analysis, and recommendation systems.
- ▶ Provides foundational embeddings for more complex models like BERT and GPT.

Contents

1 Term Frequency-Inverse Document Frequency (TF-IDF)

- Introduction to TF-IDF
- TF-IDF Example
- Advantages of TF-IDF over BoW
- Modified Versions of TF-IDF
- Hands-On: TF-IDF in Python Using Sklearn

2 Word Embeddings

- Introduction to Word Embeddings
- Limitations of Word Embeddings
- Word2Vec
- GloVe

3 Practical Application in Python

GloVe: Introduction and Core Idea

- **Developed by Stanford:**

- ▶ Created by researchers at the Stanford NLP Group, including Jeffrey Pennington, Richard Socher, and Christopher D. Manning.
- ▶ Released in 2014 to address limitations in word embeddings by integrating statistical and contextual information.

- **Core Idea:**

- ▶ Utilizes a word co-occurrence matrix, recording how frequently words appear together across a corpus.
- ▶ This matrix serves as the foundation for learning embeddings that reflect both direct and indirect word relations.

- **Training Objective:**

- ▶ Aims to reconstruct the log-probabilities of word co-occurrences, where X_{ij} is the frequency at which word j appears in the context of word i .
- ▶ Optimizes embeddings so that the dot product between two word vectors approximates the logarithm of their probability of co-occurrence.
- ▶ Captures global context from the entire corpus while also focusing on local word relationships.

GloVe: Training Objective and Advantages (Continued)

- **Advantages:**

- ▶ **Captures Complex Linguistic Patterns:**

- ★ GloVe embeddings encapsulate both semantic (meaning-related) and syntactic (grammar-related) relationships.
 - ★ Effective at distinguishing between words with similar meanings and different contexts through subtle variations captured by co-occurrence.

- ▶ **Holistic Text Understanding:**

- ★ Utilizes global statistical information by leveraging co-occurrence frequencies across the entire corpus.
 - ★ Facilitates understanding of broader language structure, supporting tasks like sentiment analysis with nuanced context comprehension.

- ▶ **Semantic Arithmetic with Embeddings:**

- ★ Supports intuitive operations, such as vector arithmetic: "king" - "man" + "woman" = "queen".
 - ★ Demonstrates the ability to perform analogy tasks, reflecting the interconnected semantic space mapped by embeddings.

Comparison: Word2Vec vs. GloVe

- **Word2Vec:**

- ▶ Uses local context.
- ▶ Faster training, lower memory footprint.

- **GloVe:**

- ▶ Combines global statistical information with local context.
- ▶ More effective for capturing deeper patterns.

- **Selecting Models:**

- ▶ Choice depends on dataset size and task specifics.
- ▶ Word2Vec for large-scale, real-time adjustments.
- ▶ GloVe for in-depth semantic understanding.

Practical Application in Python

```
from gensim.models import Word2Vec
from glove import Glove, Corpus

# Sample Data
sentences = [["cat", "sat", "on", "the", "mat"],
              ["dog", "barked"],
              ["cat", "chased", "dog"]]

# Word2Vec Model
print("Training Word2Vec model...")
word2vec_model = Word2Vec(sentences, vector_size=100, window=5, min_count=1, workers=4)
print("Word2Vec model trained.")

# Accessing Word2Vec vector
print("Word2Vec Vector for 'cat':")
print(word2vec_model.wv['cat'])

# Preparing corpus for GloVe
corpus = Corpus()
print("Fitting GloVe corpus...")
corpus.fit(sentences, window=10)

# GloVe Model
print("Training GloVe model...")
glove_model = Glove(no_components=100, learning_rate=0.05)
glove_model.fit(corpus.matrix, epochs=30, no_threads=4, verbose=True)
print("GloVe model trained.")

# Accessing GloVe vector
glove_model.add_dictionary(corpus.dictionary) # Aligning GloVe with corpus
print("GloVe Vector for 'cat':")
print(glove_model.word_vectors[corpus.dictionary['cat']])
```

Practical Application in Python

```
from gensim.models import Word2Vec
from glove import Glove, Corpus

# Sample Data
sentences = [["cat", "sat", "on", "the", "mat"],
              ["dog", "barked"],
              ["cat", "chased", "dog"]]

# Word2Vec Model
word2vec_model = Word2Vec(sentences, vector_size=100, window=5, min_count=1, workers=4)

# Preparing corpus for GloVe
corpus = Corpus()
corpus.fit(sentences, window=10)
glove_model = Glove(no_components=100, learning_rate=0.05)
glove_model.fit(corpus.matrix, epochs=30, no_threads=4, verbose=True)
```